


# KERNEL LINUX





# Introduction

# What is a kernel?

The kernel is the part of the system that:

- manages the hardware
- allocates resources eg. memory pages or CPU cycles
- is responsible for the file system and network communication
- provides an abstraction layer for the applications: the userland

# The Linux kernel

## Main features

- Portable
- Versatile
- Stable
- Mature
- Secure
- Robust

## Who is behind Linux?

- An open source community
- Around 1200 developers for one release
- Around 200 maintainers
- Around 80% of the changes are sponsored
- Linus Torvalds makes the official releases

## Usage and hardware support

Linux supports more computer architecture than any other OS.

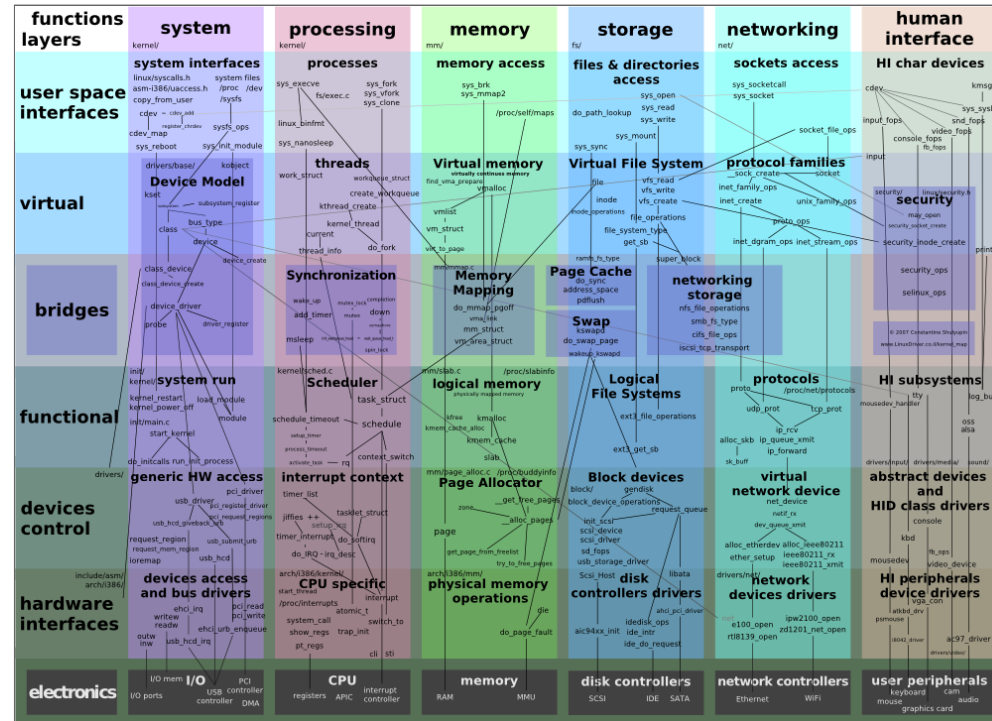
**[https://en.wikipedia.org/wiki/Usage\\_share\\_of\\_operating\\_systems](https://en.wikipedia.org/wiki/Usage_share_of_operating_systems)**

- Smartphone : 85% (android/linux)
- Desktop/Laptop : 2.7% (linux), 62% (windows)
- Server (web) : 77% (linux), 22% (windows)

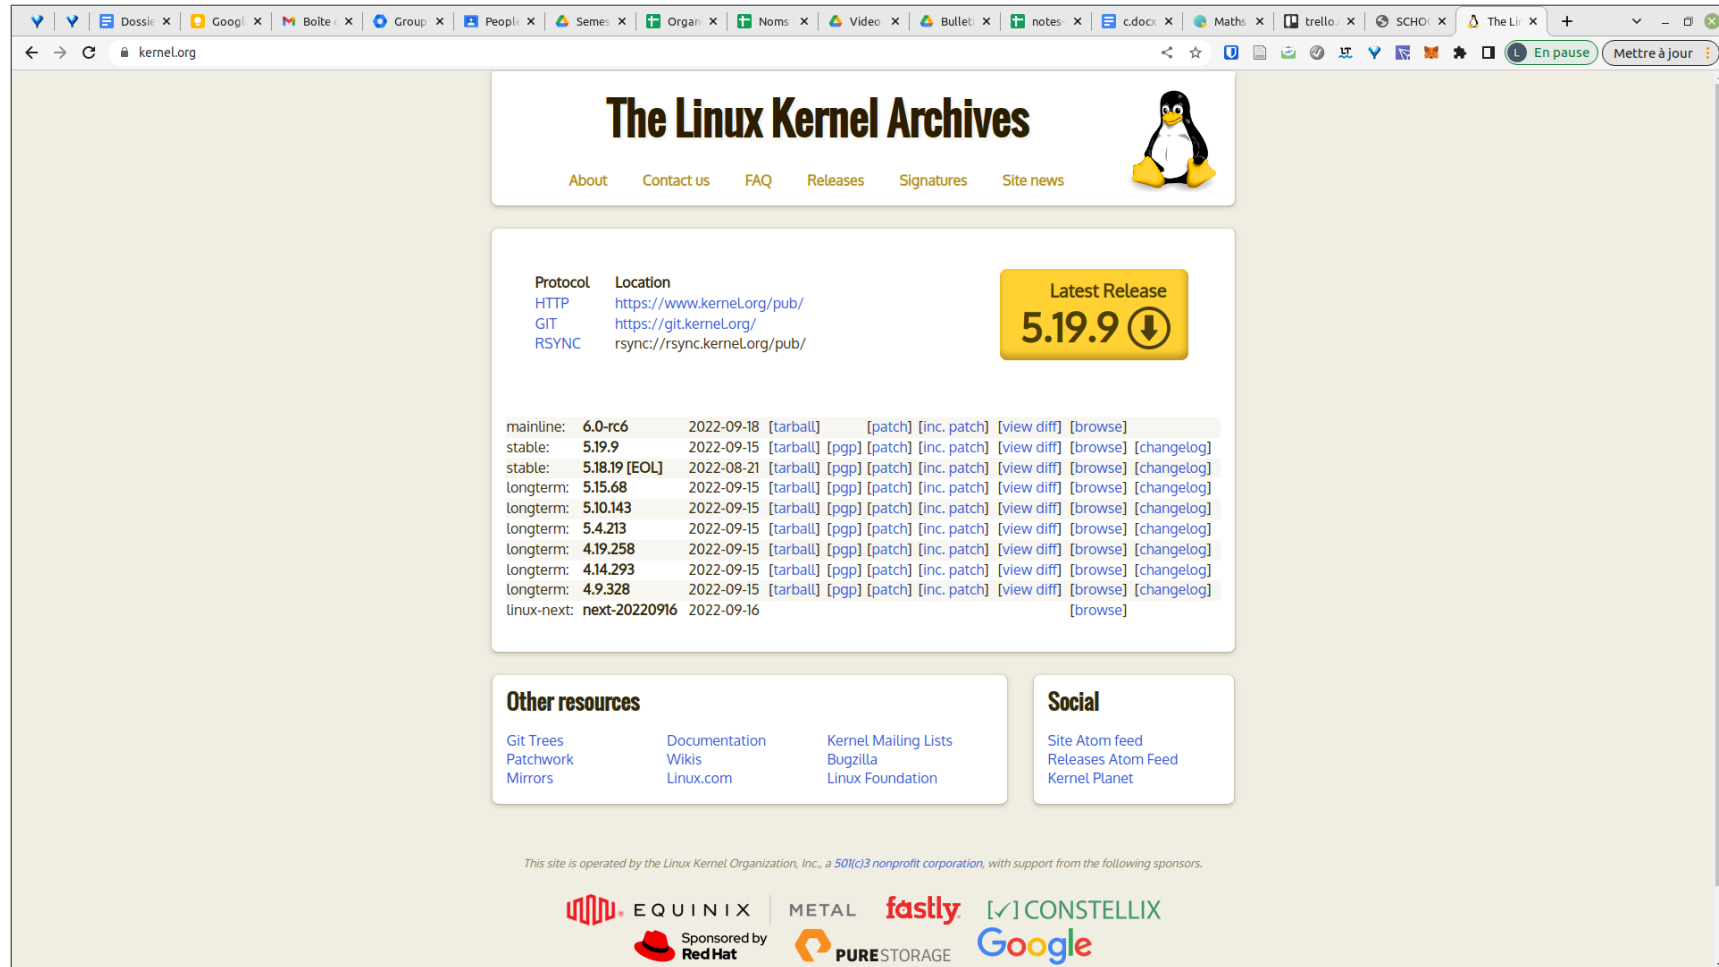


Meeting Linux

# Overview



# The source



The screenshot shows the homepage of The Linux Kernel Archives. The browser's address bar displays 'kernel.org'. The page features a navigation bar with links: About, Contact us, FAQ, Releases, Signatures, and Site news. A penguin logo is positioned to the right of the navigation bar. Below the navigation bar, there is a section for download protocols and the latest release. The protocols listed are HTTP, GIT, and RSYNC, each with its corresponding location. The latest release is 5.19.9, indicated by a yellow button with a download icon. Below this, a table lists various kernel versions, including mainline, stable, and longterm releases, along with their release dates and links to download, view diffs, and browse changelogs. At the bottom, there are sections for 'Other resources' and 'Social' links. The footer contains a disclaimer and logos of sponsors: Equinix, Metal, Fastly, Constellix, Red Hat, Pure Storage, and Google.

**The Linux Kernel Archives**

About Contact us FAQ Releases Signatures Site news

Protocol Location

HTTP <https://www.kernel.org/pub/>

GIT <https://git.kernel.org/>

RSYNC <rsync://rsync.kernel.org/pub/>

Latest Release

5.19.9

mainline:	6.0-rc6	2022-09-18	[tarball]	[patch]	[inc. patch]	[view diff]	[browse]
stable:	5.19.9	2022-09-15	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
stable:	5.18.19 [EOL]	2022-08-21	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
longterm:	5.15.68	2022-09-15	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
longterm:	5.10.143	2022-09-15	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
longterm:	5.4.213	2022-09-15	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
longterm:	4.19.258	2022-09-15	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
longterm:	4.14.293	2022-09-15	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
longterm:	4.9.328	2022-09-15	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff] [browse] [changelog]
linux-next:	next-20220916	2022-09-16					[browse]

**Other resources**

Git Trees  
Patchwork  
Mirrors


Documentation  
Wikis  
Linux.com


Kernel Mailing Lists  
Bugzilla  
Linux Foundation


**Social**


Site Atom feed  
Releases Atom Feed  
Kernel Planet


This site is operated by the Linux Kernel Organization, Inc., a 501(c)3 nonprofit corporation, with support from the following sponsors.


 EQUINIX


 Sponsored by Red Hat

 METAL

 fastly

 CONSTELLIX

 PURE STORAGE

 Google

## The kernel source tree (1/3)

- *arch/* Architecture-specific code
- *block/* Block I/O layer
- *crypto/* Cryptographic API
- *Documentation/* Kernel source documentation
- *drivers/* Device drivers (except sound ones)
- *firmware/* Device firmware needed for some devices
- *fs/* Filesystems infrastructure
- ***include/* Kernel headers**
  - *include/linux/* Core kernel headers
  - *include/uapi/* User space API headers
- *init/* Kernel initialization
- *ipc/* System V InterProcess Communication (sem, shm, msgqueue)

## The kernel source tree (2/3)

- *kernel/* Linux kernel core
- *lib/* Routine library (lists, trees, string, etc.)
- *mm/* Memory management
- *net/* Network support code (not drivers!)
- *samples/* Sample code (trace, kobject, kprobes)
- *scripts/* Scripts for internal or external use
- *security/* Security frameworks
- *sound/* Sound code and drivers
- *tools/* User space tools
- *usr/* Generate the initramfs
- *virt/* Virtualization infrastructure (kvm)

## The kernel source tree (3/3)

- *COPYING* The kernel licence (GPLv2)
- *CREDITS* The people who have contributed to the kernel
- *Kbuild* Kernel build system
- *Kconfig* Description of configuration parameters
- *MAINTAINERS* Maintainer of subsystems and drivers
- *Makefile* Base kernel Makefile
- *README* Overview of the kernel
- *REPORTING-BUGS* Instructions for reporting bugs

## Development process

- Each file has a set of maintainers.
- They are responsible for triaging bugs, reviewing patches and directing changes.
- Patches are sent to mailing-lists for review.
- Once approved, they are sent to the maintainer of the subsystem.
- Subsystem maintainers review them and place them in a special branch.
- They send this branch to Linus Torvalds, that will merge it to his branch.

## Release cycle

- The merge window is open (~2 weeks).
- Release candidates (*-rc*) are published.
- No new feature is added, only bug fixes (~6-10 weeks).
- The final release is tagged by Linus Torvalds.



## Configuring the kernel

- The kernel is a single file, resulting of the compilation process.
- Compile-time options (*-D* flags) can be used to select which features are compiled-in and their settings.
- However, after the boot process, it can load module from the filesystem at runtime to extend its features. Each module is a single file.

### **Configuring the kernel is:**

- Choosing what's going into the main file, and what will be built as modules.
- - Setting various options.

## Kernel options

Options have the form ***CONFIG\_FEATURE*** and a type, eg.

- *CONFIG\_MODULES* boolean (true/false)
- *CONFIG\_INITRAMFS\_ROOT\_UID* integer
- *CONFIG\_INITRAMFS\_SOURCE* string
- *CONFIG\_MAGIC\_SYSRQ\_DEFAULT\_ENABLE* hex
- - *CONFIG\_BTRFS\_FS* tristate (true/module/false)
- Options can depend on others.

**Two types of dependencies:**

- *A select B*, enabling A enables B
- - *A depends on B*, A is not visible until B is enabled

# Kconfig

```
config BTRFS_FS
    tristate "Btrfs filesystem support"
    select CRYPTO
    select CRYPTO_CRC32C
    select ZLIB_INFLATE
    select ZLIB_DEFLATE
    select LZO_COMPRESS
    select LZO_DECOMPRESS
    select RAID6_PQ
    select XOR_BLOCKS
```

## help

Btrfs is a general purpose copy-on-write filesystem with extents, writable snapshotting, support for multiple devices and many more features focused on fault tolerance, repair and easy administration.

[...]

# Config files

## *.config*

Simple text files, *key=value* format.

Default *.config* files

- *make defconfig*: new config with default from \$ARCH supplied defconfig
- *make x86\_64\_defconfig*: request defconfig from a platform

Editing *.config* files

- *make config*: text based
- *make menuconfig*, *make nconfig*: menu/ncurses interface
- *make xconfig*, *make gconfig*: graphical interface
- *make oldconfig*: upgrade *.config* with options from the new release

## Vendor kernels

Linux distributions typically enable a lot of kernel features and drivers, most of them are compiled as modules.

Read the config file of your kernel

Require *CONFIG\_IKCONFIG\_PROC=y*:

```
$ zcat /proc/config.gz
CONFIG_64BIT=y
CONFIG_X86_64=y
CONFIG_X86=y
CONFIG_INSTRUCTION_DECODER=y
CONFIG_OUTPUT_FORMAT="elf64-x86-64"
CONFIG_ARCH_DEFCONFIG="arch/x86/configs/x86_64_defconfig"
...
```

## Build

- Only interact with the top-directory Makefile

```
$ make menuconfig # Edit .config  
$ make # Build the kernel and modules
```

## Produces:

- ...
- *vmlinux*: ELF object of the kernel, cannot be booted
- *arch/x86/boot/bzImage*: bootable compressed kernel image
- *\*\*/\*.ko*: Modules

## More targets

- *make help*: list all available targets
- *make modules*: build/rebuild modules
- *make headers\_install*: “install” headers in the local *usr/*  
    ***make modules\_install*: install to */lib/modules/KVER***
- ▪ *INSTALL\_MOD\_PATH=dir/* to select the directory



# Interacting with Linux

## The command line: kernel parameters

### It is a string for runtime configuration:

- no recompilation
- can be builtin, using the *CONFIG\_CMDLINE* option
- - can be used to pass arguments to the *init* program

### Many kernel options, examples:

- *root=/dev/sda1* the root filesystem
- *console=ttyS0* where to write kernel messages
- *debug, loglevel=7* kernel verbosity
- - *usbcore.blinkenlights=true* also available for modules
- More documentation: *Documentation/kernel-parameters.txt*

# Syscalls

- The system call is the fundamental interface between an application and the Linux kernel.
- Typically accessed via wrapper functions of the libc. The name of the wrapper function is usually the same as the name of the system call that it invokes.
- More than 320 on x86, some are architecture-specific, but most are common.
- One of the key component of Linux' maturity.

More details in the next lesson!

## The kernel log

### **The kernel stores msgs in a circular log buffer:**

- */proc/ksmsg* for raw output
- ▪ */dev/kmsg* for structured message reading
- By default the kernel log is outputted on the console, see the *console=* kernel parameter.
- The *dmesg* tool (diagnostic message) can be used to read those messages.
- When using *systemd*, *journalctl -k* displays the kernel log.

Loadable kernel modules (LKM)

## LOADING

- Require *CONFIG\_MODULES=y*
- *insmod*: Plug a *.ko* file into the kernel.
- *modprobe*: Load a module (no *.ko*) and its dependencies.

### **Both handle module parameters:**

- - *\$ insmod ./snd-intel8x0m.ko index=-2*

## UNLOADING

- Require *CONFIG\_MODULE\_UNLOAD=y*
- *rmmmod*: Unplug the module.
- *modprobe -r*: Also remove unused dependencies.

## INFO

- *lsmod*: Show the status of modules in the Linux Kernel.
- *modinfo*: Show information about a Linux Kernel module.



## PSEUDO FILE SYSTEMS

There are many pseudo file systems in Linux, here are some of them:

- *proc*: (*/proc*) processes info, raw stuff, etc.
- *sysfs*: (*/sys*) structured information about various kernel subsystems, tied to *kobjects*
- *devtmpfs*: (*/dev*) kernel populated devices nodes

More details in the next lesson!

Writing code for Linux

# Essential developer tools

- The C language (ISO C99 and gnu extensions)
- GNU Make
- Git

# LXR (<http://lxr.free-electrons.com>)



## Linux Cross Reference

Free Electrons  
Embedded Linux Experts

• [Source Navigation](#) • [Identifier Search](#) • [Fretext Search](#) •

Version: 2.0.40 2.2.26 2.4.37 3.0 3.1 3.2 3.3 3.4 3.5 3.6 3.7 3.8 3.9 3.10 3.11 3.12 3.13 3.14 3.15 3.16

### Linux/

- Documentation/
- arch/
- block/
- crypto/
- drivers/
- firmware/
- fs/
- include/
- init/
- ipc/
- kernel/
- lib/
- mm/
- net/



## LXR (<http://lxr.free-electrons.com>)

```
930 static int __ref kernel_init(void *unused)
931 {
932     int ret;
933
934     kernel_init_freeable();
935     /* need to finish all async __init code before freeing the memory */
936     async_synchronize_full();
937     free_initmem();
938     mark_rodata_ro();
939     system_state = SYSTEM_RUNNING;
940     numa_default_policy();
941
942     flush_delayed_fput();
943
944     if (ramdisk_execute_command) {
945         ret = run_init_process(ramdisk_execute_command);
946         if (!ret)
947             return 0;
948         pr_err("Failed to execute %s (error %d)\n",
949             ramdisk_execute_command, ret);
950     }
951
952     /*
953     * We try each of these until one succeeds.
954     *
955     * The Bourne shell can be used instead of init if we are
956     * trying to recover a really broken machine.
957     */
958     if (execute_command) {
959         ret = run_init_process(execute_command);
960         if (!ret)
961             return 0;
962         pr_err("Failed to execute %s (error %d). Attempting defaults...\n",
963             execute_command, ret);
964     }
```

## Kernel space

- No access to the usual libc.
- No unbreakable memory protection.
- No floating-point operations.
- Fixed-size stack.
- Preemptive and Symetric Multi-Processors (SMP) capable.  
Synchronization and concurrency are major concerns.

Your first module



## Where to put-it?

Two alternatives:

### **Inside the official kernel tree:**

- Integrated in the kernel repository
- Can be built statically
- - Supported by the community: debug and update

### **Out of tree:**

- In a directory outside the kernel source
- - Needs to be built separately

# Hello Kernel (1/3)

Let's do some *in tree* modification:

- Add an entry in linux/init/Kconfig

```
source "my_subtree/Kconfig"
```

- Add an entry in linux/init/my\_subtree/Kconfig

```
config MY_PATCH
bool "Enable my custom patch"
help
    This is just a test
```

## Hello Kernel (2/3)

- In init/Makefile :

```
obj-$(CONFIG_MY_PATCH) += my_subtree/myfile.o
```

- Let's add some code:

```
int my_2600_function(void);  
  
int my_2600_function(void)  
{  
    return 2600;  
}
```

## Hello Kernel (3/3)

- Let's patch *init/main.c*

```
// in kernel_init
static int __ref kernel_init(void *unused)
{
    int ret;

    #if defined(CONFIG_MY_PATCH)
        ret = my_2600_function();
    #endif
    ...
}
```

# Hello Module

Let's create a *hello.c* file.

```
/* hello.c */
#include <linux/init.h>
#include <linux/module.h>
#include <linux/printk.h>

static __init int hello_init(void)
{
    pr_info("Load the module!\n");
    return 0;
}

static __exit void hello_exit(void)
{
    pr_info("Unload the module.\n");
}
```

```
}
```

```
module_init(hello_init);
```

```
module_exit(hello_exit);
```

```
MODULE_LICENSE("GPL");
```

```
MODULE_DESCRIPTION("Hello module");
```

```
MODULE_AUTHOR("Lionel Auroux");
```

# Explanations

## Macros

- `__init`: The code is discarded after initialization (both for static and module compilation).
- `__exit`: The code is discarded if the module is built statically or if module unloading is disabled.

## Module setup/cleanup

```
int mod_init(void);  
void mod_exit(void);
```

# Explanations

## Module metadata

- Retrievable using *modinfo*.

### ***MODULE\_LICENSE* is important:**

- Non-free software modules have no access to GPL exported symbols.
- - Non-free modules set the proprietary taint flag on the kernel.
- Also useful: *MODULE\_VERSION*, *MODULE\_INFO*, *MODULE\_SOFTDEP*
- More information: *include/linux/module.h*



## *printk()*

- *printk()* is the kernel space version of *printf*.

### **Multiple logging levels:**

- *pr\_emerg()*, *pr\_alert()*, *pr\_crit()*, *pr\_err()*, *pr\_warn()*,  
*pr\_notice()*, *pr\_info()*, *pr\_debug()*
- Modern code uses *pr\_xxx*.
- *pr\_devel()* while you develop your code.
- *pr\_cont()* for continuing lines with no newline (*n*).
- Formats are described in *Documentation/printk-formats.txt*.

## *pr\_fmt*

You define the *pr\_fmt* macro to set the default format to all your *pr\_xxx* calls.

```
#define pr_fmt(fmt) KBUILD_MODNAME ": " fmt
#include <linux/printk.h>

...
    pr_devel("Test.\n");
```

## Compiling your module

You need to create a *Makefile*.  
In your *Makefile*.

```
# called from kernel build system: just declare our module  
obj-m := hello.o
```

Let's compile it ...

```
user $ make -C $Where_Is_My_Linux_Source/ M=`pwd` modules
```

## Dynamic debug

- Enabled if your kernel have `CONFIG_DYNAMIC_DEBUG=y`.
- Registers a control file: `/sys/kernel/debug/dynamic_debug/control`.
- Display the current configuration by reading this file.
- Enable debug calls using:

```
# debug for the module `mymodule`  
$ cd /sys/kernel/debug/dynamic_debug  
$ echo 'module mymodule +p' > control
```

- Documentation in *Documentation/dynamic-debug-howto.txt*

## Submitting patches

### **Commit:**

1. • *git commit –signoff*

### **Make a patch:**

2. • *git format-patch origin..master*

### **Check your patch:**

3. • *./scripts/checkpatch.pl 0001-my-commit.patch*

### **Send email:**

4. • *git send-email 0001-my-commit.patch*

More documentation

## More documentation

### IRC

- *irc.oftc.net #kernelnewbies*

### Web

- **<http://kernelnewbies.org>**

### Books

Warning: Linux moves fast, those books contain outdated information.

- Linux device drivers: **LDD3**
- Linux Kernel Development
- Linux System Programming
- Linux in a Nutshell
- Understanding the Linux Kernel
- The Linux Programming Interface

# What is a syscall?

User space can issue requests to the kernel in order to access its resources or perform restricted operations.

You can think of a syscall as regular function call, but where the code being called is in the kernel.

Syscalls usages:

- Manipulating files and VFS: *open, read, write, ...*
- System setup: *gettimeofday, swapon, shutdown...*
- Processes management: *clone, mmap, ...*
- Manipulating devices: *ioctl, mount, ...*
- Cryptography and security: *seccomp, getrandom, ...*
- ...



# The syscall userland interfaces

In assembly

- On x86

```
mov eax, 1 ; exit  
int 0x80 ; or sysenter
```

- Syscall number: *eax*
- Arguments: *ebx, ecx, edx, esi, edi, ebp*, then use the stack
- On x86\_64

```
mov rax, 60 ; exit  
syscall
```

- Syscall number: *rax*

- Arguments: *rdi*, *rsi*, *rdx*, *rcx*, *r8* and *r9*, no args on memory

## syscall(2)

```
#include <unistd.h>
#include <sys/syscall.h> /* for __NR_xxx */

long syscall(long number, ...);
```

- Copies the arguments and syscall number to the registers.
- Traps to kernel code.
- Sets *errno* if the syscall returns an error.

## Don't panic!

- You will learn all about that in kernel from scratch!
- You almost never use direct calls to *syscall(2)*.
- Your libc provides wrappers for most of the syscalls you need.
- Linux also abstracts all those details in kernel code.
- For a list of the Linux system calls, see *syscalls(2)*.

## vdso(7)

- Virtual Dynamically linked Shared Objects
- Small shared library (8k) that the kernel automatically maps into the address space of all user-space applications.
- Contains non privileged code and data: *gettimeofday*, *time*, *clock\_gettime*, ... (arch-depedent)
- The ELF must be dynamically linked.

### Why?

- Making system calls can be slow.
- On x86 32bit, *int 0x80* is expensive: goes through the full interrupt-handling paths in the processor's microcode as well as in the kernel.
- Even if there is a dedicated instr (*syscall*), context switching must be done.

# Context switch

A context is:

- The CPU registers (including the instruction pointer)

## **The state of a process (including threads):**

- Memory state: stack, page tables, etc.
- CPU state: registers, caches, etc.
- Process scheduler state
- - ...

## vdso in action

```
$ cat time.c
int main(int ac, char **av) {
    printf("%d\n", time(0));
}

$ gcc time.c -o time -static
$ strace -e time ./time
time(NULL)                                = 1411171041
1411171041
+++ exited with 11 +++
$ gcc time.c -o time
$ ldd ./time
        linux-vdso.so.1 (0x00007ffffe1735000)
        libc.so.6 => /usr/lib/libc.so.6 (0x00007fee5e753000)
        /lib64/ld-linux-x86-64.so.2 (0x00007fee5eb01000)
$ strace -e time ./time
```

1411171118

+++ exited with 11 +++



## Implementation

- Defining a syscall

Include *#include <linux/syscalls.h>* and use the *SYSCALL\_DEFINE*(syscall, ...) macros anywhere in Linux code.

These macros expands to:

- *SYSCALL\_METADATA*(syscall, ...) generate metadata used in the *FTRACE* tracing framework.
- *\_\_SYSCALL\_DEFINE*(syscall, ...) more function definition expansion.
- Ultimately expand to: *asm linkage long Sys\_syscall(..)*
- *asm linkage* means that arguments are on the stack.

## Example

In *kernel/signal.c*:

```
3538 SYSCALL_DEFINE0(pause)
3539 {
3540     while (!signal_pending(current)) {
3541         current->state = TASK_INTERRUPTIBLE;
3542         schedule();
3543     }
3544     return -ERESTARTNOHAND;
3545 }
```

```
#include <asm/current.h>
...
pr_debug("The process is \"%s\" (pid %i)\n",
        current->comm, current->pid);
```

```
static inline int signal_pending(struct task_struct *p)
{
    return unlikely(
        test_tsk_thread_flag(p, TIF_SIGPENDING));
}
```

## schedule()

Ask the scheduling subsystem to pick the next process to run.

- The syscalls tables

See *arch/x86/entry/syscalls/syscall\_{32,64}.tbl*.

- *syscall\_32.tbl*

```
# <number> <abi> <name> <entry point> <compat entry point>
0          i386  restart_syscall      sys_restart_syscall
1          i386  exit                  sys_exit
2          i386  fork                  sys_fork
stub32_fork
3          i386  read                  sys_read
4          i386  write                 sys_write
5          i386  open                  sys_open
```

compat\_sys\_open

6

i386

close

sys\_close

## *syscall\_64.tbl*

0	common	read	sys_read
1	common	write	sys_write
2	common	open	sys_open
3	common	close	sys_close
4	common	stat	sys_newstat
5	common	fstat	sys_newfstat
...			
16	64	ioctl	sys_ioctl
...			
514	x32	ioctl	compat_sys_ioctl

## Generation

- Kbuild calls the *syscalltbl.sh* to generate  
*arch/x86/include/generated/asm/syscalls\_{64,32}.h*
- Same with *syscallhdr.sh*

A guided tour of some syscalls



# sysinfo

*kernel/sys.c*

```
2099 SYSCALL_DEFINE1(sysinfo,  
                    struct sysinfo __user *, info)  
2100 {  
2101     struct sysinfo val;  
2102  
2103     do_sysinfo(&val);  
2104  
2105     if (copy_to_user(info, &val,  
                      sizeof(struct sysinfo)))  
2106         return -EFAULT;  
2107
```

```
2108         return 0;  
2109     }
```

## `__user`

- Used by tools such as *sparse* to statically check the use of userspace pointers.
- `# define __user __attribute__((noderef, address_space(1)))`

## copy\_to\_user

- Copy data from kernel land to user land.

**Checks that all bytes are writeable, using:**

- - *access\_ok(VERIFY\_WRITE, addr\_to, length)*

# ioctl

```
#include <sys/ioctl.h>
int ioctl(int d, unsigned long request, ...);
```

- Control devices.
  - A big mess:**
    - Request numbers encodes data.
  - - Request data is untyped (*void \**).
- See LDD3, Chapter 6: Advanced Char Driver Operations.

# clone

```
SYSCALL_DEFINE5(clone, unsigned long, clone_flags,  
                 unsigned long, newsp,  
                 int __user *, parent_tidptr,  
                 int __user *, child_tidptr,  
                 int, tls_val)  
{  
    return do_fork(clone_flags, newsp, 0, parent_tidptr,  
child_tidptr);  
}
```

# fork

```
SYSCALL_DEFINE0(fork)
{
    return do_fork(SIGCHLD, 0, 0, NULL, NULL);
}
```

# vfork

```
SYSCALL_DEFINE0(vfork)
{
    return do_fork(CLONE_VFORK | CLONE_VM | SIGCHLD, 0,
                   0, NULL, NULL);
}
```



## References

- <http://lwn.net/Articles/604287/>
- <http://lwn.net/Articles/604515/>
- <https://www.kernel.org/doc/htmldocs/kernel-hacking>
- Searchable Linux Syscall Table: <https://filippo.io/linux-syscall-table/>